

BF527: Applications in Bioinformatics

Note: Your code should follow the guidelines laid out in class, including commenting. Partial credit will be given for nonfunctional code that is logical and well commented. This assignment must be completed on your own.

HOMEWORK 5

See [Blackboard](#) for assignment and due dates

PROBLEM 5.1 (30%):

In this problem you will be writing a python script to extract information from a **CSV** (comma separated values) file. The file is called "**blast_results.csv**" and can be downloaded from Blackboard. The file contains the top 100 hits from a BLAST search. You can open up the file in Excel to view the contents:

- **Row 1:** the headers describing each field (column).
- **Column 2:** this contains the ID for each hit (subject). Note that several hits actually contain multiple IDs corresponding to redundant entries in the NCBI **nr** database. If a hit contains multiple IDs, the IDs are separated by a semi-colon (;).
- **Column 13:** this contains the bit score for each hit.

Create a script called "**parse_blast_hits.py**" or write your code in the box below. This script should do the following:

1. Counts the total number of subject IDs in the file.

Hint: there is always *at least* 1 ID per line. If there are 2 IDs, there is 1 semi-colon (;) in the subject ID field. If there are 3 IDs, there are 2 semi-colons, etc. You can use the **count** function to count the number of semi-colons in a string.

2. Calculates the average bit score of the top 100 hits (all the scores in the file).

Hint: You can use the **int(string)** or **float(string)** functions to convert the bit score, which is stored as a string in the file, to a number. You will have to store all the bit scores to be able to later calculate the mean.

Your output should look like:

Total subject IDs: 181 Average bit score: 346.53

```
In [5]: # Write your code here
```

```

import csv #import csv module
#initialize variables
count=100 #there is always at least 1 ID per line
score=0
#open file
with open('blast_results.csv','r') as file:
    reader = csv.reader(file,delimiter=',') #read cvs file
    next(reader) #skip the first line
    for row in reader:
        count += row[1].count(';') #count ';' in second column in each row and
        score += float(row[-1]) #total score
ave_score= score/100
print('Total subject IDs:',count)
print('Average bit score:',ave_score)

```

Total subject IDs: 181
Average bit score: 346.53

PROBLEM 5.2 (40%):

Protein tyrosine kinases are implicated in several forms of cancer. In this problem you will use **ClustalW** to identify the functional tyrosine kinase domain in several proteins.

Hint: the domain is about 250 residues long and is well conserved.

- **(A)** Gather the protein sequences of the following four human tyrosine kinases from the **UniProt database** (<http://www.uniprot.org/>). A simple search of "Human" plus the gene symbols (given below) will be enough to find these four proteins. Check that the entry names you select make sense - the first search hit may not be the right one!
 1. JAK2
 2. SRC
 3. EGFR
 4. LYN

Hint: entries with **Star** are manually annotated and reviewed.

```

>sp|O60674|JAK2_HUMAN Tyrosine-protein kinase JAK2 OS=Homo sapiens OX=9606 GN=JAK2 PE=1 SV=2
MGMACLTMTMEGTSTSSIQNGDISGNANSMKQIDPVLQVYLYHSLGKSEADYLTFFPSG
EYVAEEICIAASKACGITPVYHNMFALMSETERIWYPPNHVHFHIDESTRHNVLYRIRFYF
PRWYCSGSNRAYRHGISRGAEAPLLDDFVMSYLFAQWRHDFVHGWIKVPVTHETQEECLG
MAVLDMMRIAKENDQTPLAIYNSISYKTFLPKCIRAKIQDYHILTRKRIRYRFRRIQQF
SQCKATARNLKLKYLINLETLSQAFYTEKFEVKEPGSGPSGEEIFATIIITGNGGIQWSR
GKHKESETLTEQDLQLYCDFPNIIDVSIKQANQEGSNESRVVTIHKQDGKNLEIELSSLR
EALSFVSLIDGYRRLTADAHHYLCKEVAPPVLENIQSNCHGPISMDFAISKLLKAGNQT
GLYVLRCSPKDFNKYFLTFAVERENVIEYKHCLITKNENEEYNLSGTKKNFSSLKDLLNC
YQMETVRSDNIIFQFTKCCPPKPKDKSNLLVFRTNGVSDVPTSPTLQRP THMNQMVFHKI
RNEDLIFNESLGQGTFTKIFKGVRREVGDYQGQLHETEVLLKVLDAHRNYSSESFFEAASM
MSKLSHKHLVLNYGVCVCGDENILVQEFVKFGSLDTYLLKKNKNCINILWKLEVAKQLAWA
MHFLEENTLIHGNVCAKNILLIREEDRKTGNPPFIKLSDPGISITVLPKDILQERIPWVP
PECIENPKNLNLATDKWSFGTTLWEICSGGDKPLSALDSQRKLQFYEDRHQLPAPKWAEL
ANLINNCMDYEPDFRPSFRAIIRDLNSLFTPDYELLTENDMLPNMRIGALGFSGAFEDRD
PTQFEERHLKFLQQLGKGNFGSVEMCRYDPLQDNTGEVVAVKKLQHSTEEHLRDFEREIE

```


ILKSLQHDNIVKYKGVCYSAGRRNLKLIMEYLPYGSLRDYLQKHKERIDHIKLLQYTSQI
 CKGMEYLGTKRYIHRDLATRNILVENENRVKIGDFGLTKVLPQDKEYYKVKEPGESPIFW
 YAPESLTESKFSVASDVWSFGVVLVYELFTYIEKSKSPPAEFMRMIGNDKQGQMIVFHLE
 LLKNNGRLLPRPDGCPDEIYMIMTECWNNNVNQRPSFRDLALRVDQIRDNMAG >sp|P00533|EGFR_HUMAN
 Epidermal growth factor receptor OS=Homo sapiens OX=9606 GN=EGFR PE=1 SV=2
 MRPSGTAGAALLALLAALCPASRALEEKVCQGTSNKLTQLGTFEDHFLSLQRMFNNCEV
 VLGNEITYVQRNYDLSFLKTIQEVAGYVLIALNTVERIPLNLQIIRGNMYYENSYALA
 VLSNYDANKTGLKELPMRNLQEILHGAVRFSNNPALCNVESIQWRDIVSSDFLSNMSMDF
 QNHLGSCQKCDPSCPNGSCWGAGEENCQKLTKIICAQQCSGRCRGKSPSDCCHNQCAAGC
 TGPRESDECLVCRKFRDEATCKDTCPLMLYNPTTYQMDVNPEGKYSFGATCVKKCPRNYV
 VTDHGSCVRACGADSYEMEEDGVRKCKKCEGPCRKVCNGIGIGEFKDSLSINATNIKHFK
 NCTSIGDLHLIPVAFRGDSFTHTPPLDPQELDILKTVKEITGFLLIQAWPENRTDLHAF
 ENLEIIRGRTKQHGQFSLAVVSLNITSLGLRSLKEISDGDVIISGNKNLCYANTINWKKL
 FGTSGQKTKIISNRGENSCKATGQVCHALCSPEGCWGPEPRDCVSCRNVSRGRECVDKCN
 LLEGEPREFVENSECIQCHPECLPQAMNITCTGRGPDNCIQCAHYIDGPHCVKTCAPAGVM
 GENNTLVWKYADAGHVCHLCHPNCTYGCTGPGLEGCPPTNGPKIPSIATGMVGALLLLLVV
 ALGIGLFMRRRHIVRKRTLRRLLQERELVEPLTPSGEAPNQALLRILKETEFKKIKVLGS
 GAFGTVYKGLWIPEGEKVKIPVAIKELREATSPKANKEILDEAYVMASVDNPHVCRLGI
 CLTSTVQLITQLMPFGCLLDYVREHKDNIGSQYLLNWCVQIAKGMNYLEDRLVHRDLAA
 RNVLVKTPQHVKITDFGLAKLLGAEKEYHAEGGKVPIKWMALESILHRIYTHQSDVWSY
 GVTVWELMTFGSKPYDGIPASEISSILEKGERLPQPPICTIDVYMIMVKCWMIDADSRPK
 FRELIIEFSKMARDPQRYLVIQGDERMHLPSPTDSNFYRALMDEEDMDDVDDADEYLIPQ
 QGFFSSPSTSRTPLSSLATSNNSTVACIDRNLQSCPIKEDSFLQRYSSDPTGALTED
 SIDDTFLPVPEYINQSVPKRPAGSVQNPVYHNQPLNPAPSRDPHYQDPHSTAVGNPEYLN
 TVQPTCVNSTFDSPAHWAKGSHQISLDNPDYQQDFFPKEAKPNGIFKGSTAENAEYLRV APQSSEFIGA
 >sp|P07948|LYN_HUMAN Tyrosine-protein kinase Lyn OS=Homo sapiens OX=9606 GN=LYN PE=1 SV=3
 MGCIKSKGKDSLSDDGVDLKTQPVNRTERTIYVRDPTSNKQQRVPESQLLPGQRFQTKD
 PEEQGDIVVALYPYDGIHPDDL SFKKGEKMKVLEEHGEWWKAKSLLTKKEGFIPSNYVAK
 LNTLETEEWFFKDITRKDAERQLLAPGNSAGAFIRESETLKGSFSLSVRDFDPVHGDVI
 KHYKIRSLDNGGYYISPRITFPCISDMIKHYQKQADGLCRRLEKACISPKPQKPWDKDAW
 EIPRESIKLVKRLGAGQFGEVWMGYNNSTKVAVKTLKPGTMSVQAFLEEANLMKTLQHD
 KLVRLYAVVTTREEPIYITEYMAKGSLLDFLKSDEGGKVLLPKLIDFSAQIAEGMAYIER
 KNYIHRDLRAANVLVSESLMCKIADFGLARVIEDNEYTAREGAKFPIKWTAPEAINFGCF
 TIKSDVWSFGILLYEIVTYGKIPYPGRTNADVMTALSQGYRMPRVENCPELYDIMKMCW
 KEKAEERPTFDYLQSVLDDFYTATEGYQQQP >sp|P12931|SRC_HUMAN Proto-oncogene tyrosine-protein
 kinase Src OS=Homo sapiens OX=9606 GN=SRC PE=1 SV=3
 MGSNKSXPKDASQRRRSLEPAENVHGAGGGAFFASQTPSKPASADGHRGPSAAFAPAAAE
 PKLFGGFNSSDVTVTSPQRAGPLAGGVTTFVALYDYESRTETDLSFKKGERLQIVNNTGDD
 WWLAHSLSTGQTGYIPSNYVAPSDSIQAEWYFGKITRRESERLLLNAENPRGTFLVRES
 ETTKGAYCLSVSDFDNAKGLNVKHYKIRKLDSGGFYITSRTQFNSLQQLVAYYSKHADGL
 CHRLTTVCPTSKPQTQGLAKDAWEIPRESLRLEVKLGGQCFGEVWMGTWNGTTRVAIKTL
 KPGTMSPEAFLQEAQVMKKLRHEKLVQLYAVVSEPIYIVTEYMSKGSLLDFLKGETGKY
 LRLPQLVDMAAQIASGMAYVERMNYVHRDLRAANILVGENLVCKVADFGLARLIEDNEYT
 ARQGAKFPIKWTAPEAALYGRFTIKSDVWSFGILLTELTTKGRVPYPGMVNREVLDQVER
 GYRMPCPPECPESLHDLMCQCWRKEPEERPTFEYLQAFLEDYFTSTEPQYQPGENL

- **(B)** Use ClustalW to align the four protein sequences. Qualitatively and quantitatively evaluate the alignment, i.e., does this look like a good alignment? Does the alignment score support your opinion?

Yes, this looks like a good alignment as there're a lot conserved residues and regions. And the highest pairwise score is 52.3438, multiple sequence alignment score is 2472.

- **(C)** Identify the tyrosine kinase domain. Specifically, report its start and stop positions in the alignment. Provide a screenshot of the JalView output for part of the domain.

The tyrosine kinase domain starts near 750 and stops near 1000 in the alignment.



- **(D)** Replace one of the tyrosine kinase sequences with an unrelated protein sequence of your choice. Report the sequence that you used. Rebuild the alignment and compare it to the one obtained in **(B)**.

I replaced JAK2 with IL-7(Interleukin-7). >sp|P13232|IL7_HUMAN Interleukin-7 OS=Homo sapiens OX=9606 GN=IL7 PE=1 SV=1 MFHVSFRYIFGLPPLILVLLPVASSDCDIEGKDGKQYESVLMVSIQQLLDSDMKEIGSNCL NNEFNFFKRHICDANKEGMFLFRAARKLRQFLKMNSTGDFDLHLLKVSEGTILLNCTGQ VKGRKPAALGEAQPTKSLEENKSLKEQKKLNDLCFLKRLLEIKTCWNKILMGTKEH

- **(E)** Can you still identify the tyrosine kinase domain even though you have thrown an unrelated sequence into the mix? Why or why not?

Yes, but with lower certainty. Because we have multiple sequences, the other three tyrosine kinase proteins can still provide a good alignment in that tyrosine kinase domain, although the unrelated sequences introduce some disturbances.

PROBLEM 5.3 (30%):

Thought question: describe in **English words** or **pseudocode** how you would write a program that works like a virtual ribosome, i.e. a script that takes an mRNA sequence and translates it to its corresponding protein sequence. Assume that the input mRNA sequence

is stored in a FASTA file and the output protein sequence must be written out to a FASTA file.

Hint: your program will need some additional information to be able to translate from mRNA to protein, which you must describe how to store and use.

First, establish a library of genetic code (mRNA codons as keys, amino acids as values, e.g. {"UUU":"F","UUC":"F"}, there're 3 stop codons with a value "STOP"); Input, read and save mRNA sequence FASTA file as mRNA; Create a new FASTA file to store the output protein sequence; For i in the length of mRNA: If mRNA[i:i+3] is "AUG" ("start" signal to kick off translation): translation starts, add 'M' to protein; read every three nucleotides from here and find corresponding amino acids from genetic code library defined at the beginning; add amino acid to the protein sequence one by one; If met a stop codon: stop translation. End End End Close mRNA file. Close protein file.